

Logistic Regression

Classification

Regression v/s Classification

- *In regression, we predict continuous values.*
- *In classification, we predict categories or labels.*

For example:

- *A typical regression problem would be something like, predict price of a house given some parameters.*
- *Whereas, a classification problem would be something like, predict if this tumor is malignant or benign given some parameters.*

Logistic Regression

Definition

- *Logistic regression is a classification algorithm that predicts the probability of an event happening.*
- *It is mainly used for binary classification (e.g., Yes/No , Pass/Fail).*

Working

- *We predict a probability between 0 and 1.*
- *Eg) If probability > 0.5 , classify as 1 (Yes/True).
If probability < 0.5 , classify as 0 (No/False).*

The sigmoid function


What is it?

- Its a mathematical function that converts any number into a value between 0 and 1

Mathematical denotion:

$$g(z) = \frac{1}{1+e^{-z}} \quad 0 < g(z) < 1$$

Where: $z=wx+b$



Cost function

- We cant simply use MSE for classification because its cost curve is not convex in this case
- Hence, we formulate the following loss function

$$L(f_{\vec{w},b}(\vec{x}^{(i)}), y^{(i)}) = \begin{cases} -\log(f_{\vec{w},b}(\vec{x}^{(i)})) & \text{if } y^{(i)} = 1 \\ -\log(1 - f_{\vec{w},b}(\vec{x}^{(i)})) & \text{if } y^{(i)} = 0 \end{cases}$$

Simplified form:

$$L(f_{\vec{w},b}(\vec{x}^{(i)}), y^{(i)}) = -y^{(i)}\log(f_{\vec{w},b}(\vec{x}^{(i)})) - (1 - y^{(i)})\log(1 - f_{\vec{w},b}(\vec{x}^{(i)}))$$

Hence, we write cost function as:

$$\overset{\text{loss}}{L(f_{\vec{w},b}(\vec{x}^{(i)}), y^{(i)})} = -y^{(i)}\log(f_{\vec{w},b}(\vec{x}^{(i)})) - (1 - y^{(i)})\log(1 - f_{\vec{w},b}(\vec{x}^{(i)}))$$

$$\overset{\text{cost}}{J(\vec{w}, b)} = \frac{1}{m} \sum_{i=1}^m [L(f_{\vec{w},b}(\vec{x}^{(i)}), y^{(i)})]$$

Gradient Decent

Just like during linear regression, we now use gradient decent to minimize the cost function

$$w_j = w_j - \alpha \frac{\partial}{\partial w_j} J(\vec{w}, b)$$

$$b = b - \alpha \frac{\partial}{\partial b} J(\vec{w}, b)$$


The partial derivatives look like this:

$$\frac{\partial}{\partial w_j} J(\vec{w}, b) = \frac{1}{m} \sum_{i=1}^m (f_{\vec{w}, b}(\vec{x}^{(i)}) - y^{(i)}) x_j^{(i)}$$

$$\frac{\partial}{\partial b} J(\vec{w}, b) = \frac{1}{m} \sum_{i=1}^m (f_{\vec{w}, b}(\vec{x}^{(i)}) - y^{(i)})$$


Decision boundary

Lets revisit the sigmoid curve:



Here, to classify, we set a threshold

- If $\sigma(z) \geq 0.5 \rightarrow$ predict class 1.
- If $\sigma(z) < 0.5 \rightarrow$ predict class 0.



So, we have defined the decision boundary at: $wX+b=0$

Performance Metrics

Now that we've trained our model, we need to see how well it has performed.

Accuracy is not a good measure for classification as it can be misleading.

Consider this case, if a dataset has 98% +ve results and only 2% -ve results. A model that only predicts $y=1$ will have 98% accuracy!

How to solve this? Enter, Confusion Matrix: From this , we can calculate Precision and Recall:

		Actual Values	
		Positive (1)	Negative (0)
Predicted Values	Positive (1)	TP	FP
	Negative (0)	FN	TN

$$\text{Precision} = \frac{\text{True Positive}}{\text{True Positive} + \text{False Positive}}$$

$$\text{Recall} = \frac{\text{True Positive}}{\text{True Positive} + \text{False Negative}}$$

Performance Metrics

It is often convenient to combine precision and recall into a single metric called the F_1 score. It's just harmonic mean of Precision and Recall

$$F1 \text{ score} = \frac{2}{\frac{1}{\text{Precision}} + \frac{1}{\text{Recall}}} = 2 \cdot \frac{\text{Precision} * \text{Recall}}{\text{Precision} + \text{Recall}}$$

$$\Rightarrow F1 \text{ score} = 2 \cdot \frac{\text{Precision} * \text{Recall}}{\text{Precision} + \text{Recall}}$$

Precision Recall Tradeoff: Generally as you increase threshold, your precision increases and recall decreases and vice versa. Based on what you want to do, you may want higher precision or higher recall.

